

INTRODUCTION

- The purpose of the lab is to use naive bayes to get probabilities for text classification
 - The primary objective is to classify texts and to accurately predict the section role (BACKGROUND, METHODS, RESULTS, OBJECTIVE, CONCLUSION) of biomedical abstract sentences
 - **PART-A: Multinomial Naive Bayes (MNB)**
 - **PART-B: TF-IDF score based classifier: Term Frequency-Inverse Document Frequency**
 - **PART-C: BAYES OPTIMAL CLASSIFIER**
-

METHODOLOGY

MULTINOMIAL NAIVE BIAS

- Multinomial Naive Bayes (MNB) is a variant of Naive Bayes designed specifically for discrete count data.
- We use log prior $\log P(C)$ and log likelihood $\log P(w_i|C)$ along with laplace(Additive) smoothening where $\alpha = 1$.

BAYES OPTIMAL CLASSIFIER

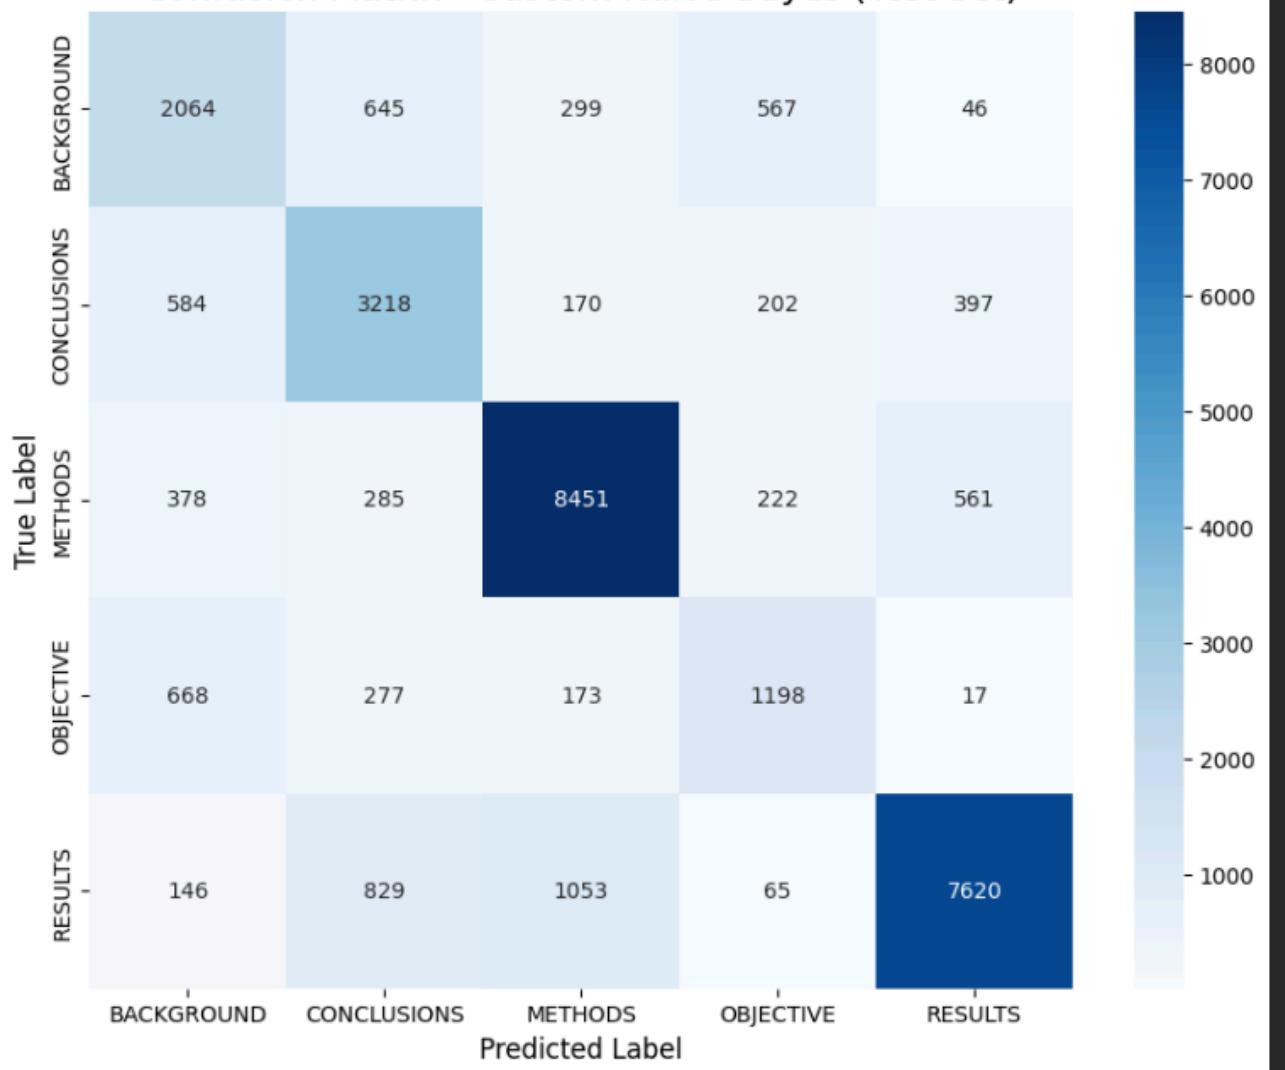
- The **Bayes Optimal Classifier** represents the absolute best possible classifier for a given problem
 - 5 models were chosen :Multinomial Naive Bayes, Logistic Regression, Random Forest, Decision Tree, and K-Nearest Neighbors
 - Posterior probability for each hypothesis was calculated and The sampled training data (`x_train_sampled`) was first split into a sub-training set (80%) and a validation set (20%).
 - All models were trained on this sample
-

RESULT AND ANALYSIS

PART-A

Generating Confusion Matrix...

Confusion Matrix - Custom Naive Bayes (Test Set)



== Test Set Evaluation (Custom Count-Based Naive Bayes) ==

Accuracy: 0.7483

	precision	recall	f1-score	support
BACKGROUND	0.54	0.57	0.55	3621
CONCLUSIONS	0.61	0.70	0.66	4571
METHODS	0.83	0.85	0.84	9897
OBJECTIVE	0.53	0.51	0.52	2333
RESULTS	0.88	0.78	0.83	9713
accuracy			0.75	30135
macro avg	0.68	0.69	0.68	30135
weighted avg	0.76	0.75	0.75	30135

Macro-averaged F1 score: 0.6809

PART-B

```
→ Training initial Naive Bayes pipeline...
Training complete.

== Test Set Evaluation (Initial Sklearn Model) ==
Accuracy: 0.7266
      precision    recall  f1-score   support

 BACKGROUND       0.64     0.43     0.51     3621
 CONCLUSIONS      0.62     0.61     0.62     4571
 METHODS          0.72     0.90     0.80     9897
 OBJECTIVE         0.73     0.10     0.18     2333
 RESULTS           0.80     0.87     0.83     9713

 accuracy           -       -       0.73     30135
 macro avg        0.70     0.58     0.59     30135
 weighted avg      0.72     0.73     0.70     30135

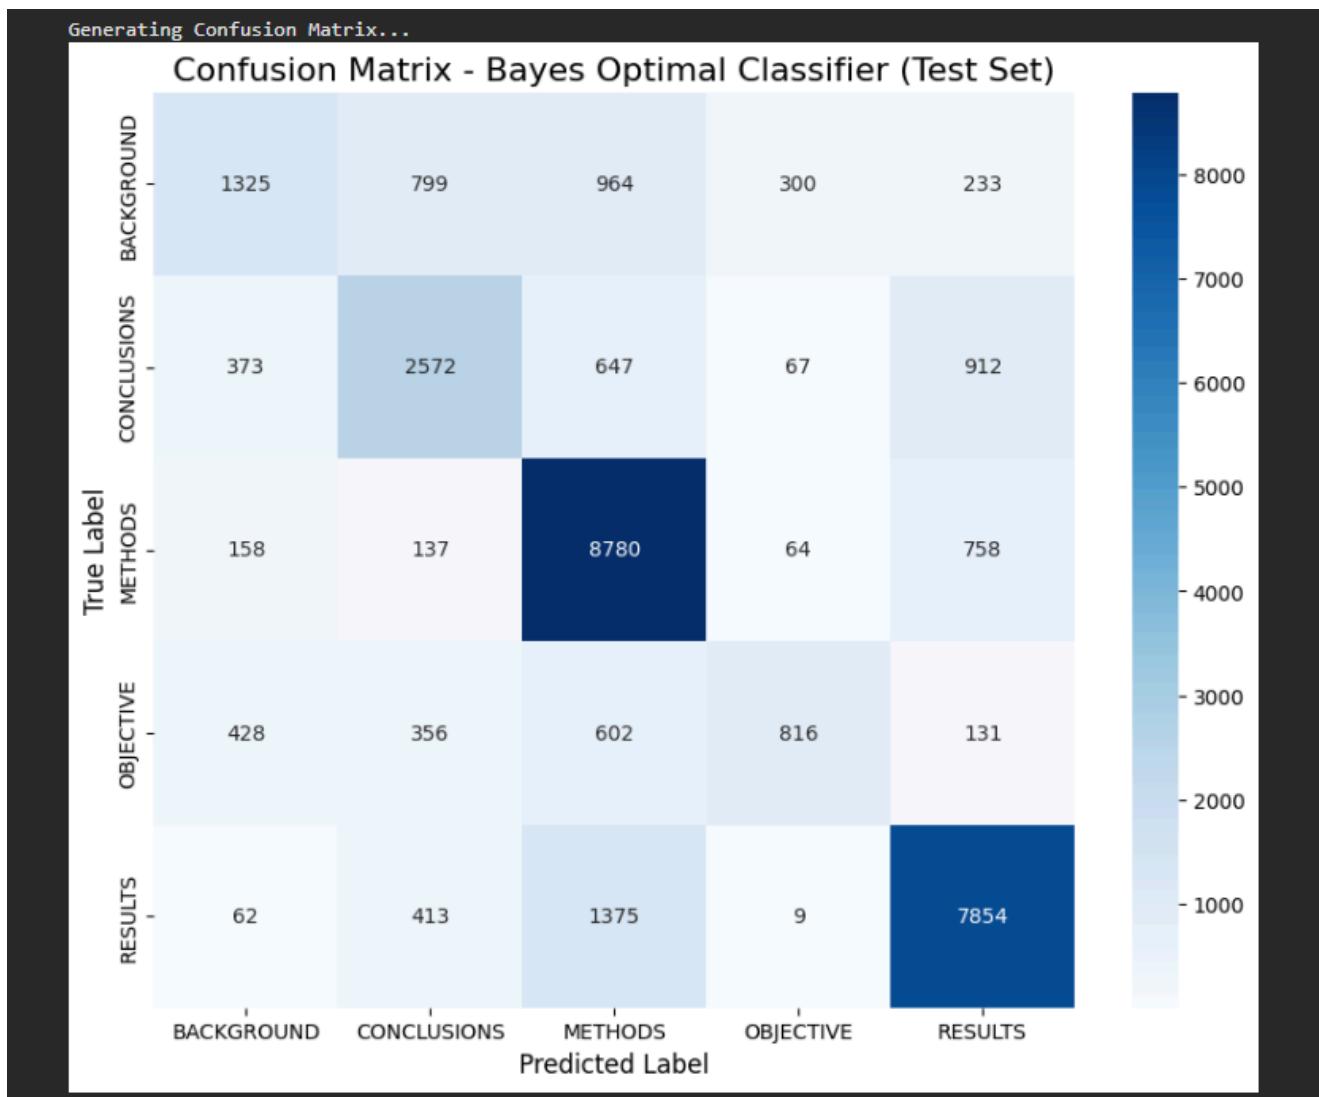
Macro-averaged F1 score: 0.5877

Starting Hyperparameter Tuning on Development Set...
Fitting 3 folds for each of 6 candidates, totalling 18 fits
Grid search complete.

Best Parameters (from Dev Set): {'nb_alpha': 0.1, 'tfidf_ngram_range': (1, 2)}
Best CV F1-macro score (from Dev Set): 0.6567
```

PART-C

```
→ Please enter your full SRN (e.g., PES1UG22CS345): PES2UG23CS148
Using dynamic sample size: 10148
Actual sampled training set size used: 10148
```



```

LogisticRegression Log-Likelihood: -1833.8303
Fitting RandomForest on sub-train split...
RandomForest Log-Likelihood: -2076.4980
Fitting DecisionTree on sub-train split...
DecisionTree Log-Likelihood: -2533.2165
Fitting KNN on sub-train split...
KNN Log-Likelihood: -2945.7222

Calculated Posterior Weights ( $P(h|D)$ ):
NaiveBayes: 0.000000
LogisticRegression: 1.000000
RandomForest: 0.000000
DecisionTree: 0.000000
KNN: 0.000000

Fitting the VotingClassifier (BOC approximation)...
Fitting complete.

Predicting on test set...

== Final Evaluation: Bayes Optimal Classifier (Soft Voting) ==
Accuracy: 0.7084
Macro-averaged F1 score: 0.6141

Classification Report:
precision    recall    f1-score   support
BACKGROUND    0.56     0.37     0.44      3621
CONCLUSIONS   0.60     0.56     0.58      4571
METHODS       0.71     0.89     0.79      9897
OBJECTIVE     0.65     0.35     0.45      2333
RESULTS        0.79     0.81     0.80      9713

accuracy          0.71      30135
macro avg       0.66     0.59     0.61      30135
weighted avg    0.70     0.71     0.69      30135

```

DISCUSSION

My scratch model from Part A worked, but its F1 score wasn't as high as the sklearn models. This is probably because I used CountVectorizer (simple counts) and a default alpha=1.0 .

The Part B sklearn model performed better after tuning. The GridSearchCV found the best alpha and ngram_range , and using TfidfVectorizer likely made a big difference.

The Part C BOC model was the most interesting. By combining five different types of models, it didn't have to rely on just one of them being perfect. Using soft voting with posterior weights let the ensemble trust the more confident models for a given prediction. This combined approach seems more robust than just relying on a single, highly-tuned algorithm.